

CS 751: Assignment 2

Bharath Kongara

Spring 2015

Contents

1	Question 1	2
	1.1 Solution	2
2	Question 2	9
	2.1 Solution	9

1 Question 1

Choose 100 URIs from Assignment1 and generate WARC files of those URIs using:

- wget
- WARCcreate
- Heritrix (stand-alone or via WAIL)
- and webrecorder.io

Describe the resulting WARC files: quantitatively compare and contrast the results of the WARC files of the same URI as generated by different tools

- choose interesting examples

Demonstrate playback of 2-3 WARCs in the (Wayback Machine (via WAIL or stand-alone) or pywb) and (webrecorder.io)

- “<https://github.com/iipc/openwayback>”
- “<https://github.com/ikreymer/pywb>”

1.1 Solution

The following steps were taken to setup tools and generate WARC files:

- Installed wget using ‘brew install wget’.
- WARC file for each URI is generated using the command ‘wget -warc-file=outputfilename URI’.
- Downloaded WARCcreate chrome extension from chrome web store and added it to the chrome browser. Generated WARC for each URI manually by providing input to WARCcreate.
- Installed WAIL.
- The figure below shows how to generate WARC files for multiple URIs using single heritrix instance.

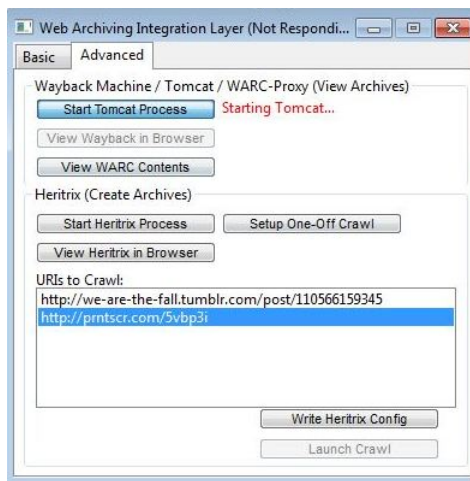


Figure 1: Generating WARC files for multiple URIs using WAIL

- Of the 100 URI selected, when any of the URIs had a parameter then WAIL wasn't able to generate WARC file for it.
- Used “<https://webrecorder.io/>” . Generating WARC files for multiple URIs is done as shown in the figure below.

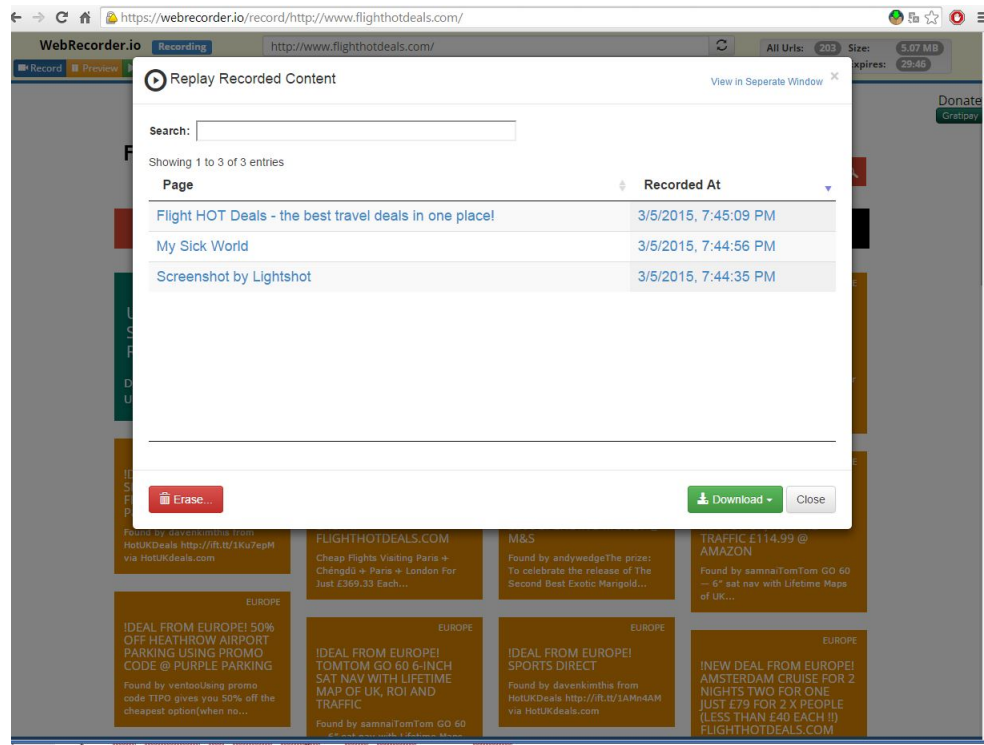


Figure 2: Generating WARC files for multiple URIs using webrecorder

- Quantitative comparison of WARC files generated by WAIL, webrecorder, WARCcreate and wget is done on the basis of WARC file sizes.
- The comparison sizes are WAIL = 40 MB , webrecorder.io = 20.48 MB, WARCcreate = 15.36 MB, wget = 4 MB.
- The WAIL size is more compared to others because WAIL software crawls data of the links in the website.
- This comparison is shown in the graph below.
- Installed pywb to playback WARC files.
- pywb requires cdx for each WARC to playback, so generated .cdx file for each WARC file.
- I have put all my WARC files into a new folder and changed the archive paths of config.xml to point to my WARC files.
- Play back for two WARC files using pywb is shown in the below figure.
- The WARC file which was in the play back does not have some images and some URLs.
- The WARC file which is in the play back does not contain certain CSS archives and some links
- WebRecorder.io takes WARC file to replay the archived file. Play back for two WARC files using WebRecorder.io is shown in the below figure
- The WARC file in this play back worked perfectly without any faults.
- The WARC file in this play back did not have certain images.

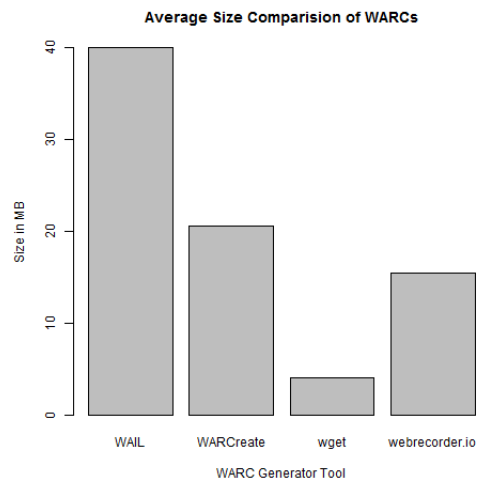


Figure 3: Average Size Comparison of WARC



Figure 4: Play back WARC file using pywb

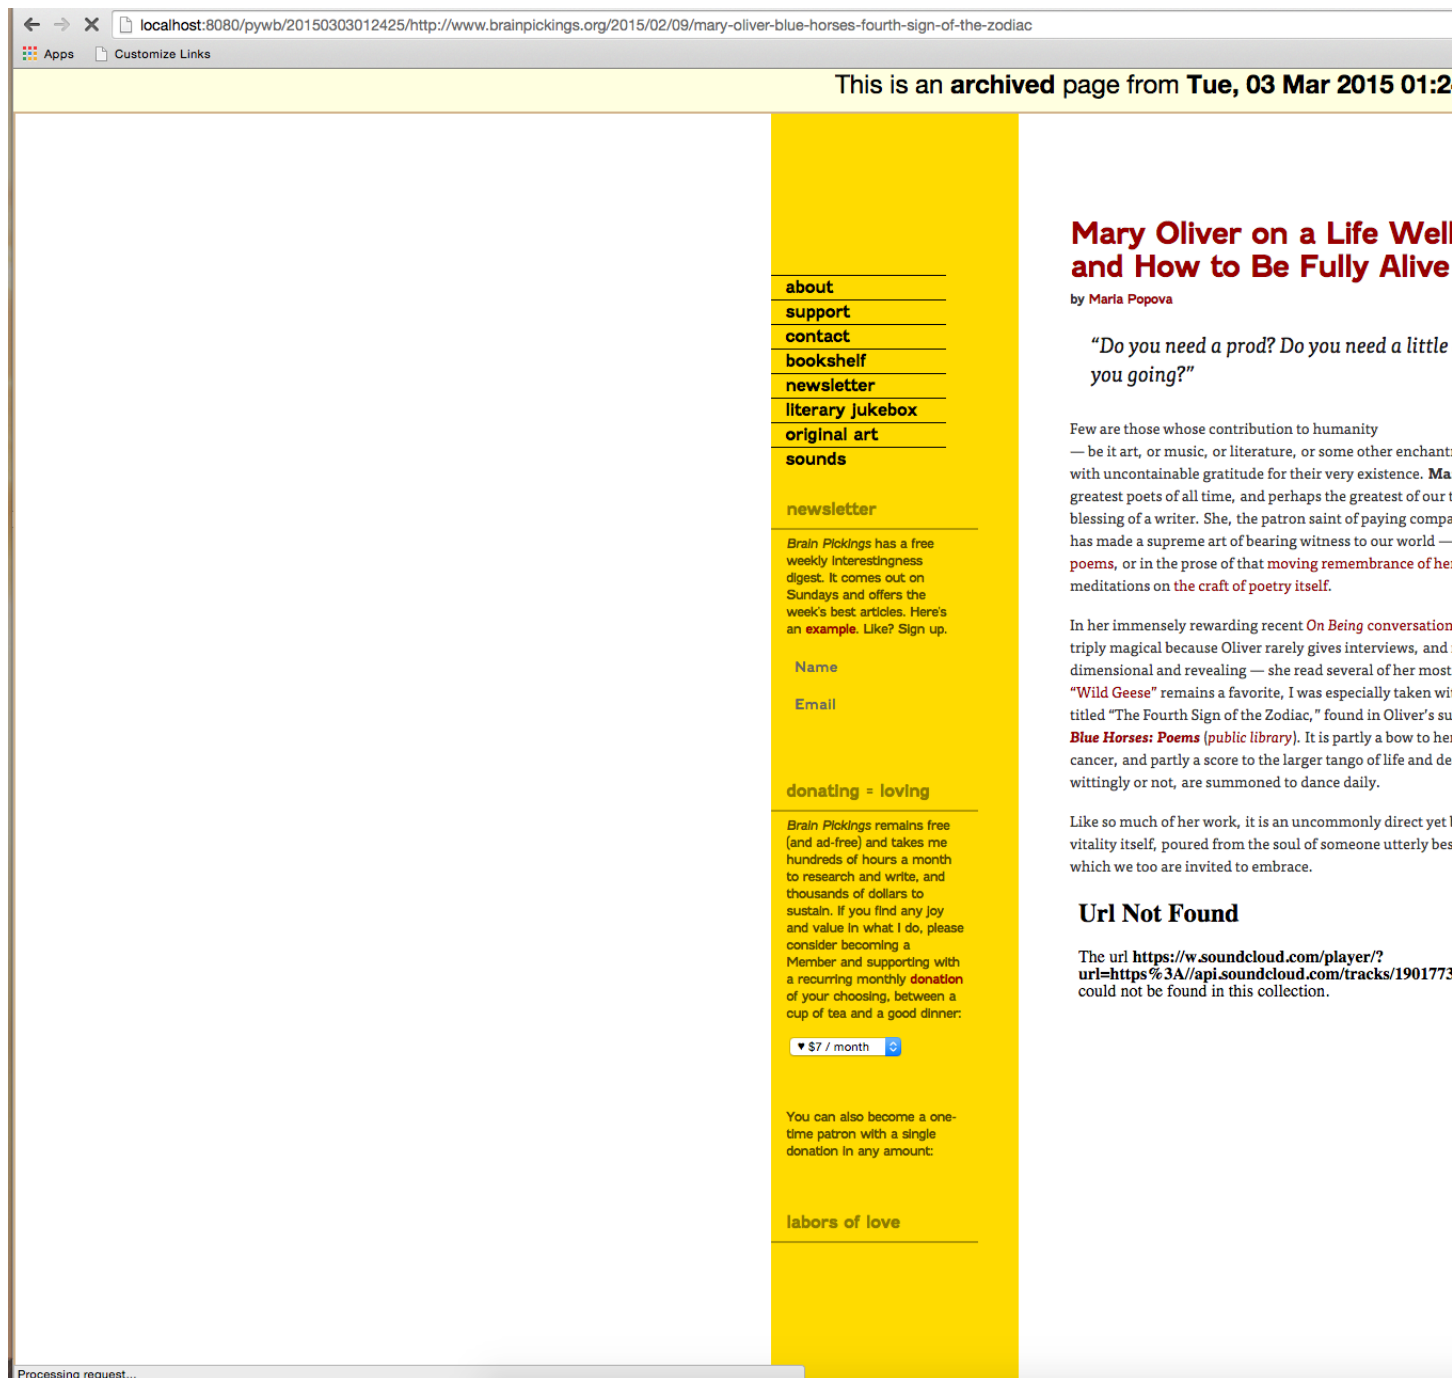


Figure 5: Play back WARC file using pywb



Figure 6: Play back WARC file using webrecorder

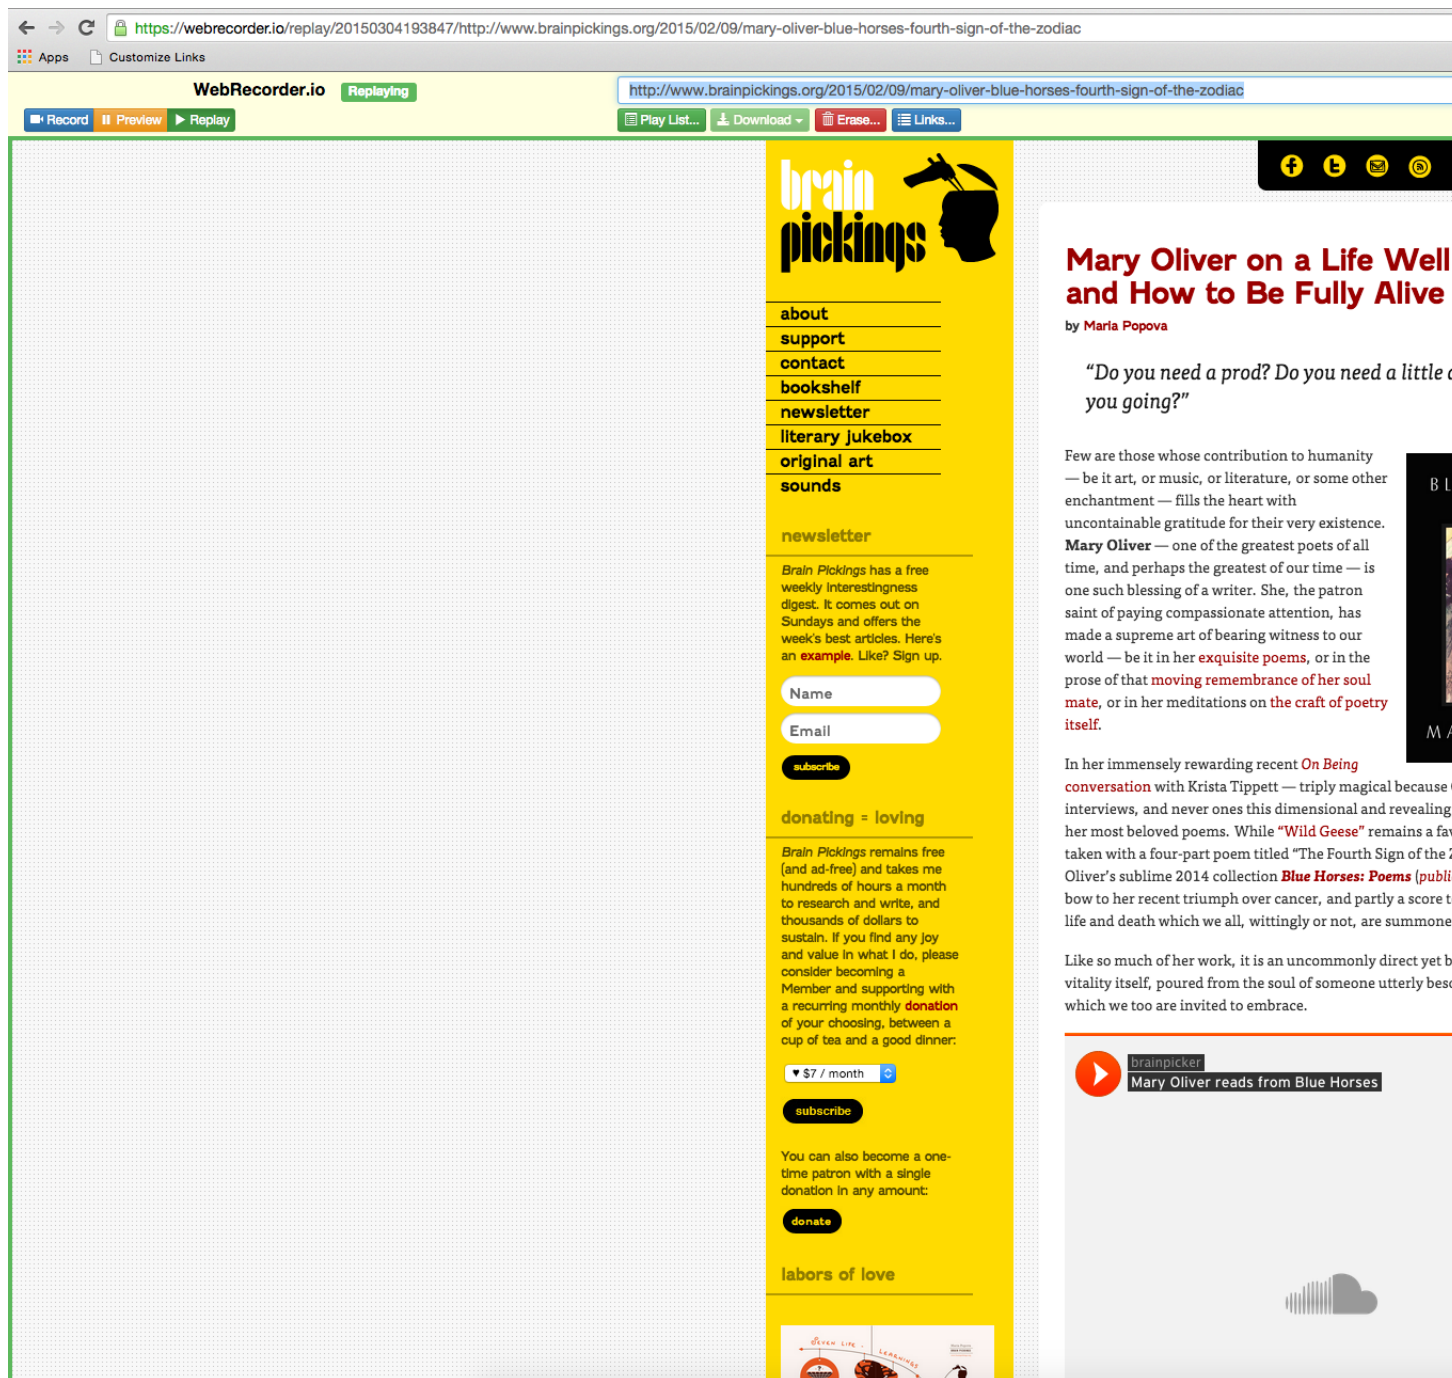


Figure 7: Play back WARC file using webrecorder

2 Question 2

- Ingest the 100 URIs from their resulting WARC files into a SOLR instance - see the code + tutorial at: <https://github.com/ukwa/webarchive-discovery>
- Demonstrate several functioning queries on the files(a full front-end is not required) - describe the configuration choices you made in setting up SOLR and processing the documents

2.1 Solution

The following steps were taken to configure SOLR and process documents :

- The pre-requisites for SOLR are Maven 3, Java 7 so I installed them.
- I faced an issue while installing SOLR which is jetty dependency not found.
- I got it working by adding the dependency to the pom.xml
- The command “mvn jetty:run-exploded” starts the SOLR instance.
- I indexed the WARC file by using the command “java -jar path of jar -s “<http://localhost:8080/discovery>” -t Path of WARC”.
- I tried to set-up Shine for SOLR as the front end instead of the default UI for SOLR but after trying for two days and multiple email exchanges with the author of the tool I wasn’t able to. So I used the default SOLR front end.
- In SOLR we can perform the following queries.
 - Here we are demonstrating how to retrieve the names and ids of all documents with “<http://localhost:8080/solr/select?q=inStock:false&wt=json&fl=id,name>”
 - Here we are using the functional query idf(field,term). This function returns the inverse document frequency for the given term, using the similarity for the field. “[http://localhost:8080/solr/select/?fl=score,id&defType=func&q=mul\(tf\(text,memory\),idf\(text,memory\)\)](http://localhost:8080/solr/select/?fl=score,id&defType=func&q=mul(tf(text,memory),idf(text,memory)))”
 - The functional query being used here is tf(field,term) and it returns the inverse document frequency factor for the given term using the similarity for the field. “[http://localhost:8080/solr/select/?fl=score,id&defType=func&q=mul\(tf\(\\$f,\\$t\),idf\(\\$f,\\$t\)\)&f=text&t=memory](http://localhost:8080/solr/select/?fl=score,id&defType=func&q=mul(tf($f,$t),idf($f,$t))&f=text&t=memory)”
 - termfreq(field,term) is the functional query that is being used and it returns the number of times the term term appers for that field in the document. “[http://localhost:8080/solr/select/?fl=score,id&q=DDR&sort=termfreq\(text,memory\)desc](http://localhost:8080/solr/select/?fl=score,id&q=DDR&sort=termfreq(text,memory)desc)”
 - Here norm(field) is the functional query that is being used. It returns the norm stored in the index, the product of the index time boost and then length normalization factor. “[http://localhost:8983/solr/select/?fl=score,id&q=DDR&sort=norm\(text\)asc](http://localhost:8983/solr/select/?fl=score,id&q=DDR&sort=norm(text)asc)”.
- Below are the snapshots of the queries that I ran on SOLR after indexing.

Request-Handler (qt)

/select

— common —

q
archive

fq

sort

start.rows
0 10

fl

df

Raw Query Parameters
key1=val1&key2=val2

wt
json

☒ indent

☐ debugQuery

☐ dismax

☐ edismax

☐ hl

☐ facet

☐ spatial

☐ spellcheck

Execute Query

```
http://localhost:8080/discovery/select?q=archive&wt=json&indent=true
```

```
{
  "responseHeader": {
    "status": 0,
    "QTime": 3,
    "params": {
      "indent": "true",
      "q": "archive",
      "_": "1425699246639",
      "wt": "json"
    }
  },
  "response": {
    "numFound": 1,
    "start": 0,
    "docs": [
      {
        "source_file_s": "WARCMerge20150305163111072566.warc@1372691",
        "url": "http://www.scientificamerican.com/article/an-electrode-in-the-brain-turns-off-depression/?utm_source=feedburner",
        "host": "www.scientificamerican.com",
        "domain": "scientificamerican.com",
        "public_suffix": "com",
        "server": [
          "Apache/2.2.15 (Red Hat)"
        ],
        "content_type_served": "text/html;charset=UTF-8",
        "content_length": 54885,
        "id": "sha1:RD1W0ICPUG2K4V16FPBJRU35DZ4UNMNP/K+fgtaymSUGShX8v/K4KHw==",
        "hash": [
          "sha1:RD1W0ICPUG2K4V16FPBJRU35DZ4UNMNP"
        ],
        "crawl_date": "2015-03-03T19:57:00Z",
        "crawl_year": "2015",
        "wayback_date": "20150303195700",
        "content": [
          "0006000ADVERTISEMENTSIn | Register0Subscription Center6 Print Issues + 6 Digital Issues + 1 Year Archive6 Digital Issues Tabl"
        ],
        "content_text_length": 5844,
        "content_type": [
          "text/html"
        ],
        "title": "An Electrode in the Brain Turns Off Depression - Scientific American",
        "keywords": "electrode, mood disorder, brain, electrical stimulation, depression",
        "content_encoding": "UTF-8",
      }
    ]
  }
}
```

Figure 8: Query 1 on SOLR

Request-Handler (qt)

/select

— common —

q
india

fq

sort

start.rows
0 10

fl

df

Raw Query Parameters
key1=val1&key2=val2

wt
json

☒ indent

☐ debugQuery

☐ dismax

☐ edismax

☐ hl

☐ facet

☐ spatial

☐ spellcheck

Execute Query

```
http://localhost:8080/discovery/select?q=india&wt=json&indent=true
```

```
{
  "responseHeader": {
    "status": 0,
    "QTime": 1,
    "params": {
      "indent": "true",
      "q": "india",
      "_": "1425699315752",
      "wt": "json"
    }
  },
  "response": {
    "numFound": 1,
    "start": 0,
    "docs": [
      {
        "source_file_s": "WARCMerge20150305163111072566.warc@5870199",
        "url": "http://english.manoramaonline.com/home.html/",
        "host": "english.manoramaonline.com",
        "domain": "manoramaonline.com",
        "public_suffix": "com",
        "content_type_served": "text/html;charset=UTF-8",
        "server": [
          "Apache/2.2.15 (CentOS)"
        ],
        "content_length": 178941,
        "id": "sha1:NWFF0DWJLQ5VASFEHSDQYXTQAZQ4GC0/K6zPZXExambT9eNXX+nvkg==",
        "hash": [
          "sha1:NWFF0DWJLQ5VASFEHSDQYXTQAZQ4GC0"
        ],
        "crawl_date": "2015-03-03T19:56:58Z",
        "crawl_year": "2015",
        "wayback_date": "20150303195658",
        "content": [
          "00009CABMake Us Your Home PageFollow Us OnThiruvananthapuram20CMistAdvanced SearchAdvanced SearchEnter word or phrase"
        ],
        "content_text_length": 13856,
        "content_type": [
          "text/html"
        ],
        "title": "Home",
        "keywords": "",
        "content_encoding": "UTF-8",
      }
    ]
  }
}
```

Figure 9: Query 2 on SOLR

- Dashboard
- Logging
- Core Admin
- Java Properties
- Thread Dump
- discovery
- Overview
- Analysis
- Dataimport
- Documents
- Files
- Ping
- Plugins / Stats
- Query
- Replication
- Schema Browser

Request-Handler (qt)

/select

common

q

radio, university

fq

sort

start, rows

0 10

fl

df

Raw Query Parameters

key1=val1&key2=val2

wt

json

☒ Indent

☐ debugQuery

☐ dismax
☐ edismax
☐ hl
☐ facet
☐ spatial
☐ spellcheck

Execute Query

http://localhost:8080/discovery/select?q=radio%2C+university&wt=json&indent=true

```

{
  "responseHeader": {
    "status": 0,
    "QTime": 0,
    "params": {
      "indent": "true",
      "q": "radio, university",
      "_": "1425699461150",
      "wt": "json"
    }
  },
  "response": {
    "numFound": 1,
    "start": 0,
    "docs": [
      {
        "source_file_s": "WARCHerge20150305163111072566.warc@390479",
        "url": "http://www.toneradio.co.uk/",
        "url_type": "slashpage",
        "host": "www.toneradio.co.uk",
        "domain": "toneradio.co.uk",
        "public_suffix": "co.uk",
        "content_type_served": "text/html; charset=UTF-8",
        "server": [
          "PHP/5.5.20",
          "PleskLin",
          "cloudflare-nginx"
        ],
        "content_length": 28241,
        "id": "sha1:QXMP4RE7J2M4EBHUMI3UIYBEAB2TX0Z5/mf2c0nmvE2tp9Qof0wEy9w==",
        "hash": [
          "sha1:QXMP4RE7J2M4EBHUMI3UIYBEAB2TX0Z5"
        ],
        "crawl_date": "2015-03-03T19:54:15Z",
        "crawl_year": "2015",
        "wayback_date": "20150303195415",
        "content": [
          "6311f64Tone RadioStudent Radio for the University of GloucestershireOn-AirThe Sam & Joe ShowMessageListenM
        ],
        "content_text_length": 2264,
        "content_type": {
          "text/html"
        }
      }
    ]
  }
}

```

Figure 10: Query 3 on SOLR

Bibliography

- [1] Download warcreate. <https://chrome.google.com/webstore/detail/warcreate/kenncghfghgolcbmckhiljgaabnpcaaa?hl=en-US>.
- [2] Generating simple bar plots in r. <http://www.statmethods.net/graphs/bar.html>.
- [3] Playback using webrecorder. <https://webrecorder.io/>.
- [4] Setting up pywb. <https://github.com/ikreymer/pywb>.
- [5] Using wail. <http://matkelly.com/wail/>.
- [6] Warc merge. <https://github.com/maturban/WARCMerge>.